

# Extending Theano

Arnaud Bergeron

February 1, 2017

# Outline

1. How to Make an Op (Python) (45 min)
2. How to Make an Op (C) (30 min)
3. Op Params (10 min)
4. Optimizations (20 min)

# How to Make an Op (Python)

# Overview

```
from theano import Op

class MyOp(Op):
    __props__ = ()

    def __init__(self, ...):
        # set up parameters

    def make_node(self, ...):
        # create apply node

    def perform(self, node, inputs, outputs_storage):
        # do the computation
```

`--init--`

```
def __init__(self, ...):  
    # set up parameters
```

- ▶ Optional, a lot of Ops don't have one
- ▶ Serves to set up Op-level parameters
- ▶ Should also perform validation on those parameters

`--props--`

```
--props-- = ()
```


- ▶ Optional (although very useful)
- ▶ Generates `__hash__`, `__eq__` and `__str__` methods if present
- ▶ Empty tuple signifies no properties that should take part in comparison
- ▶ If you have only one property, make sure you add a final comma: `('property',)`

## make\_node

```
def make_node(self, ...):  
    # create apply node
```

- ▶ This creates the node object that represents our computation in the graph
- ▶ The parameters are usually Theano variables, but can be python objects too
- ▶ The return value must be an `Apply` instance

## What Is an Apply Node?





perform

```
def perform(self, node, inputs, outputs_storage):  
    # do the computation
```

- ▶ This performs the computation on a set of values (hence the method name)
- ▶ The parameters are all python objects (not symbolic values)
- ▶ This method must not return its result, but rather store it in the 1-element lists (or cells) provided in `outputs_storage`
- ▶ The output storage may contain a pre-existing value from a previous run that may be reused for storage.

# DoubleOp

```
from theano import Op, Apply
from theano.tensor import as_tensor_variable

class DoubleOp(Op):
    __props__ = ()

    def make_node(self, x):
        x = as_tensor_variable(x)
        return Apply(self, [x], [x.type()])

    def perform(self, node, inputs, output_storage):
        x = inputs[0]
        z = output_storage[0]
        z[0] = x * 2
```

## Op Instances and Nodes

When you call an op class you get an instance of that Op:

```
double_op = DoubleOp()
```

But when you want to use that op as a node in a graph you need to call the *instance*:

```
node = double_op(x)
```

You can do both steps at once with a double call like this:

```
node = DoubleOp()(x)
```

## Basic Tests

```
import numpy

from theano import function, config
from theano.tensor import matrix
from theano.tests import unittest_tools as utt
from doubleop import DoubleOp

def test_doubleop():
    utt.seed_rng()
    x = matrix()
    f = function([x], DoubleOp()(x))
    inp = numpy.asarray(numpy.random.rand(5, 4),
                        dtype=config.floatX)

    out = f(inp)
    utt.assert_allclose(inp * 2, out)
```

## Run Tests

The simplest way to run your tests is to use `nosetests` directly on your test file like this:

```
$ nosetests test_doubleop.py  
.
```

---

```
Ran 1 test in 0.427s
```

OK

You can also use `theano-nose` which is a wrapper around `nosetests` with some extra options.

## infer\_shape

```
def infer_shape(self, node, input_shapes):  
    # return output shapes
```

- ▶ This function is optional, although highly recommended
- ▶ It takes as input the symbolic shapes of the input variables
- ▶ `input_shapes` is of the form  
[[`i0_shp0`, `i0_shp1`, ...], ...]
- ▶ It must return a list with the symbolic shape of the output variables

## Example

```
def infer_shape(self, node, input_shapes):  
    return input_shapes
```

- ▶ Here the code is really simple since we don't change the shape in any way in our Op
- ▶ `input_shapes` would be an expression equivalent to `[x.shape]`

# Tests

```
from theano.tests import unittest_tools as utt

class test_Double(utt.InferShapeTester):
    def test_infer_shape(self):
        utt.seed_rng()
        x = matrix()
        self._compile_and_check(
            # function inputs (symbolic)
            [x],
            # Op instance
            [DoubleOp()(x)],
            # numeric input
            [numpy.asarray(numpy.random.rand(5, 4),
                           dtype=config.floatX)],
            # Op class that should disappear
            DoubleOp)
```



# Gradient

```
def L_op(self, inputs, outputs, output_grads):  
    # return gradient graph for each input
```

- ▶ This function is required for graphs including your op to work with `theano.grad()`
- ▶ Each item you return represents the gradient with respect to that input computed based on the gradient with respect to the outputs (which you get in `output_grads`).
- ▶ It must return a list of symbolic graphs for each of your inputs
- ▶ Inputs that have no valid gradient should have a special `DisconnectedType` value

## Example

```
def L_op(self, inputs, outputs, output_grads):  
    return [output_grads[0] * 2]
```

- ▶ Here since the operation is simple the gradient is simple
- ▶ Note that we return a list

# Tests

To test the gradient we use `verify_grad`

```
from theano.tests import unittest_tools as utt

def test_doubleop_grad():
    utt.seed_rng()
    utt.verify_grad(
        # Op instance
        DoubleOp(),
        # Numeric inputs
        [numpy.random.rand(5, 7, 2)]
    )
```

It will compute the gradient numerically and symbolically (using our `L_op()` method) and compare the two.

# How to Make an Op (C)

# Overview

```
from theano import Op

class MyOp(Op):
    __props__ = ()

    def make_node(self, ...):
        # return apply node

    def c_code(self, node, name, input_names,
               output_names, sub):
        # return C code string

    def c_support_code(self):
        # return C code string

    def c_code_cache_version(self):
        # return hashable object
```

## c\_code

```
def c_code(self, node, name, input_names,  
           output_names, sub):  
    # return C code string
```

- ▶ This method returns a python string containing C code
- ▶ `input_names` contains the variable names where the inputs are
- ▶ `output_names` contains the variable names where to place the outputs
- ▶ `sub` contains some code snippets to insert into our code (mostly to indicate failure)
- ▶ The variables in `output_names` may contain a reference to a pre-existing value from a previous run that may be reused for storage.

## Support Code

```
def c_support_code(self):  
    # return C code string
```

- ▶ This method return a python string containing C code
- ▶ The code may be shared with multiple instances of the op
- ▶ It can contain things like helper functions

There are a number of similar methods to insert code at various points

## Headers, Libraries, Compilers

Some of the methods available to customize the compilation environment:

`c_libraries` Return a list of shared libraries the op needs

`c_headers` Return a list of included headers the op needs

`c_compiler` C compiler to use (if not the default)

Again others are available. Refer to the documentation for a complete list.



## Python C-API

- `void Py_INCREF(PyObject *o)` Increase the reference count of a python object.
- `void Py_DECREF(PyObject *o)` Decrease the reference count of a python object.
- `void Py_XINCREF(PyObject *o)` Increase the reference count of a (potentially NULL) python object.
- `void Py_XDECREF(PyObject *o)` Decrease the reference count of a (potentially NULL) python object.

## Numpy C-API

`int PyArray_NDIM(PyArrayObject *a)` Get the number of dimension of an array.

`numpy_intp *PyArray_DIMS(PyArrayObject *a)` Get the shape of an array.

`numpy_intp *PyArray_STRIDES(PyArrayObject *a)` Get the strides of an array.

`void * PyArray_DATA(PyArrayObject *a)` Get the data pointer (pointer to element 0) of an array.

## Example I

This is the C code equivalent to `perform`

```
from theano import Op, Apply
from theano.tensor import as_tensor_variable

class DoubleC(Op):
    __props__ = ()

    def make_node(self, x):
        x = as_tensor_variable(x)
        if x.ndim != 1:
            raise TypeError("DoubleC only works on 1D")
        return Apply(self, [x], [x.type()])
```

## Example II

```
def c_code(self, node, name, input_names,
            output_names, sub):
    return """
Py_XDECREF(%(out)s);
%(out)s = (PyArrayObject *)PyArray_NewLikeArray(
    %(inp)s, NPY_ANYORDER, NULL, 0);
if (%(out)s == NULL) {
    %(fail)s
}
for (npyp_intp i = 0; i < PyArray_DIM(%(inp)s, 0); i++) {
    *(dtype_%(out)s *)PyArray_GETPTR1(%(out)s, i) =
        (*(dtype_%(inp)s *)PyArray_GETPTR1(%(inp)s, i)) * 2;
}
""" % dict(inp=input_names[0], out=output_names[0],
           fail=sub["fail"])
```

# COp

```
from theano.gof import COp

class MyOp(COp):
    __props__ = ()

    def __init__(self, ...):
        COp.__init__(self, c_files, func_name)
        # Other init code if needed

    def make_node(self, ...):
        # make the Apply node
```

# Constructor Arguments

- ▶ Basically you just pass arguments to the constructor of COp
  - ▶ Either by calling the constructor directly  
`COp.__init__(self, ...)`
  - ▶ Or via the superclass **super** (`MyOp, self`) `__init__(...)`
- ▶ The arguments are:
  - ▶ a list of file names with code sections (relative to the location of the op class)
  - ▶ the name of a function to call to make the computation (optional)

## COp: Example

```
from theano import Apply
from theano.gof import COp
from theano.tensor import as_tensor_variable

class DoubleCOp(COp):
    __props__ = ()

    def __init__(self):
        COp.__init__(self, ["doublecop.c"],
                      "APPLY_SPECIFIC(doublecop)")

    def make_node(self, x):
        x = as_tensor_variable(x)
        if x.ndim != 1:
            raise TypeError("DoubleCOp only works with 1D")
        return Apply(self, [x], [x.type()])
```

## COp: Example

```
#section support_code

int APPLY_SPECIFIC(doublecop) (PyArrayObject *x,
                               PyArrayObject **out) {
    Py_XDECREF(*out);
    *out = (PyArrayObject *)PyArray_NewLikeArray(
        inp, NPY_ANYORDER, NULL, 0);

    if (*out == NULL)
        return -1;

    for (npy_intp i = 0; i < PyArray_DIM(x, 0); i++) {
        *(DTYPE_OUTPUT_0 *)PyArray_GETPTR1(*out, i) =
            (*(DTYPE_INPUT_0 *)PyArray_GETPTR1(x, i)) * 2;
    }
    return 0;
}
```



# Tests

- ▶ Testing ops with C code is done the same way as testing for python ops
- ▶ One thing to watch for is tests for ops which don't have python code
  - ▶ You should skip the test in those cases
  - ▶ Test for `theano.config.gxx == ""`
- ▶ Using `DebugMode` will compare the output of the Python version to the output of the C version and raise an error if they don't match

## Gradient and Other Concerns

- ▶ The code for `grad()` and `infer_shape()` is done the same way as for a python Op
- ▶ In fact you can have the same Op with a python and a C version sharing the `grad()` and `infer_shape()` code
  - ▶ That's how most Ops are implemented

# Op Params

# Purpose

- ▶ Used to pass information to the C code
- ▶ Can reduce the amount of compiled C code
- ▶ Required for things that can change from one script run to the other.

## Usage

```
from theano import Op

class MyOp(Op):
    params_type = # a params type here

    def __init__(self, ...):
        # Get some params

    # signature change
    def perform(self, node, inputs, out_storage, params):
        # do something

    def get_params(self, node):
        # Return a params object
```

# GPU Ops

# Overview

```
from theano import Op
from theano.gpuarray.type import gpu_context_type

class GpuOp(Op):
    __props__ = ()
    params_type = gpu_context_type

    def make_node(self, ...):
        # return apply node

    def get_params(self, node):
        return node.outputs[0].type.context
```

# Overview

```
def perform(self, node, inputs, output_storage):  
    # python code  
  
def c_code(self, node, name, input_names,  
           output_names, sub):  
    # return C code string
```

- ▶ `params_type` is new.
- ▶ `get_params` is new.



## Context and Context Name

- ▶ Context is what is used to refer to the chosen GPU.  
It is a C object that can't be serialized.
- ▶ Context Name is a name internal to Theano to refer to a given context object. It is a python string.
- ▶ Context Names are used whenever you need a symbolic object.

## Double on GPU

```
try:
    from pygpu import gpuarray
except ImportError:
    pass

class DoubleGpu(Op, GpuKernelBase):
    __props__ = ()

    def make_node(self, x):
        ctx_name = infer_context_name(x)
        x = as_gpuarray_variable(x, ctx_name)
        return Apply(self, [x], [x.type()])

    def get_params(self, node):
        return node.outputs[0].type.context
```

## Double on GPU

```
def gpu_kernels(self, node, name):
    dt = node.inputs[0].type
    code = """
KERNEL void double(GLOBALMEM %(ctype) *out,
                   GLOBALMEM const %(ctype)s *a,
                   ga_size n) {
    for (ga_size i = LID_0; i < n; i += LDIM_0) {
        out[i] = 2 * a[i];
    }
}
""" % dict(ctype=gpuarray.dtype_to_ctype(dt))
    return [Kernel(code=code, name="double",
                   params=[gpuarray.GpuArray,
                           gpuarray.GpuArray,
                           gpuarray.SIZE],
                   flags=Kernel.get_flags(dt))]
```

## Double on GPU

```
def c_code(self, node, name, inn, outn, sub):
    return """
size_t n = 1;
Py_XDECREF(%(out)s);
%(out)s = pygpu_empty(PyGpuArray_NDIM(%(inp)s),
                      PyGpuArray_DIMS(%(inp)s),
                      GA_C_ORDER, %(ctx)s, Py_None);
if (%(out)s == NULL) %(fail)s
for (unsigned int i = 0; i < %(inp)s->ga.nd; i++)
    n *= PyGpuArray_DIM(%(inp)s, i);
if (double_scall(1, &n, 0, %(out)s, %(inp)s, n)) {
    PyErr_SetString(PyExc_RuntimeError,
                    "Error calling kernel");
    %(fail)s;
}
""" % dict(inp=inn[0], out=outn[0], fail=sub["fail"])
```

## GpuKernelBase

```
class DoubleCGpu(CGpuKernelBase):  
    __props__ = ()  
  
    def __init__(self):  
        CGpuKernelBase.__init__(self, ["doublecgpu.c"],  
                                   "double_fn")  
  
    def make_node(self, x):  
        ctx_name = infer_context_name(x)  
        x = as_gpuarray_variable(x, ctx_name)  
        return Apply(self, [x], [x.type()])  
  
    def get_params(self, node):  
        return node.outputs[0].type.context
```

## GpuKernelBase

```
#section kernels
#kernel double : *, *, size :

KERNEL void double(GLOBAL_MEM DTYPE_o0 *out,
                  GLOBAL_MEM DTYPE_i1 *a,
                  ga_size n) {
    for (ga_size i = LID_0; i < n; i += LDIM_0) {
        out[i] = 2 * a[i];
    }
}
```

## GpuKernelBase

```
#section support_code_struct
int double_fn(PyGpuArrayObject *inp, npy_int64 n,
              PyGpuArrayObject **out,
              PyGpuContextObject *ctx) {
    size_t n = 1;
    Py_XDECREF(*out);
    *out = pygpu_empty(PyGpuArray_NDIM(inp),
                      PyGpuArray_DIMS(inp),
                      GA_C_ORDER, ctx, Py_None);
    if (*out == NULL) return -1;
    for (unsigned int i = 0; i < inp->ga.nd; i++)
        n *= PyGpuArray_DIM(inp, i);
    if (double_scall(1, &n, 0, *out, inp, n)) {
        PyErr_SetString(PyExc_RuntimeError,
                        "Error calling kernel");
    }
    return -1;
}
```

# Optimizations



# Purpose

- ▶ End goal is to make code run faster
- ▶ Sometimes they look after stability or memory usage
- ▶ Most of the time you will make one to insert a new Op you wrote

## Replace an Op

Here is code to use `DoubleOp()` instead of `ScalMul(2)`.

```
from scalmulop import ScalMulV1
from doubleop import DoubleOp
from theano.gof import local_optimizer
from theano.tensor.opt import register_specialize

@register_specialize
@local_optimizer([ScalMulV1])
def local_scalmul_double(node):
    if not (isinstance(node.op, ScalMulV1) and
            node.op.scal == 2):
        return False

    return [DoubleOp()(node.inputs[0])]
```

## Replace an Op for GPU

Here is code to move the Double op to GPU.

```
from scalmulop import ScalMulV1
from doubleop import DoubleOp
from doublecop import DoubleCOp
from doublec import DoubleC
from doublecgpu import DoubleCGpu
from theano.gpuarray.opt import (register_opt, op_lifter,
                                  register_opt2)

@register_opt('fast_compile')
@op_lifter([DoubleOp, DoubleC, DoubleCOp])
@register_opt2([DoubleOp, DoubleC, DoubleCOp],
                'fast_compile')
def local_scalmul_double_gpu(op, context_name, inputs,
                              outputs):
    return DoubleCGpu
```

# Tests

```
import theano

from scalmulop import ScalMulV1
from doubleop import DoubleOp
import opt

def test_scalmul_double():
    x = theano.tensor.matrix()
    y = ScalMulV1(2)(x)
    f = theano.function([x], y)

    assert not any(isinstance(n.op, ScalMulV1)
                     for n in f.maker.fgraph.toposort())
    assert any(isinstance(n.op, DoubleOp)
                for n in f.maker.fgraph.toposort())
```

## Exercise

- ▶ Implement a `ScalMulOp` that multiplies its input by an arbitrary scalar value. Start with a python implementation
- ▶ Add C code to your implementation
- ▶ Create a GPU version of your op.
- ▶ Create an optimization that replace the CPU version with a GPU version when appropriate.

Clone the repo at

[https://github.com/abergeron/ccw\\_tutorial\\_theano.git](https://github.com/abergeron/ccw_tutorial_theano.git).